

ConvNetJS CIFAR-10 demo

Description

This demo trains a Convolutional Neural Network on the [CIFAR-10 dataset](#) in your browser, with nothing but Javascript. The state of the art on this dataset is about 90% accuracy and human performance is at about 94% (not perfect as the dataset can be a bit ambiguous). I used [this python script](#) to parse the [original files](#) (python version) into batches of images that can be easily loaded into page DOM with img tags.

This dataset is more difficult and it takes longer to train a network. Data augmentation includes random flipping and random image shifts by up to 2px horizontally and vertically.

By default, in this demo we're using Adadelta which is one of per-parameter adaptive step size methods, so we don't have to worry about changing learning rates or momentum over time. However, I still included the text fields for changing these if you'd like to play around with SGD+Momentum trainer.

Report questions/bugs/suggestions to [@karpathy](#).

Training Stats

pause

Forward time per example: 8ms

Backprop time per example: 13ms

Classification loss: 1.92149

L2 Weight decay loss: 0.00102

Training accuracy: 0.34

Validation accuracy: 0.1

Examples seen: 1002

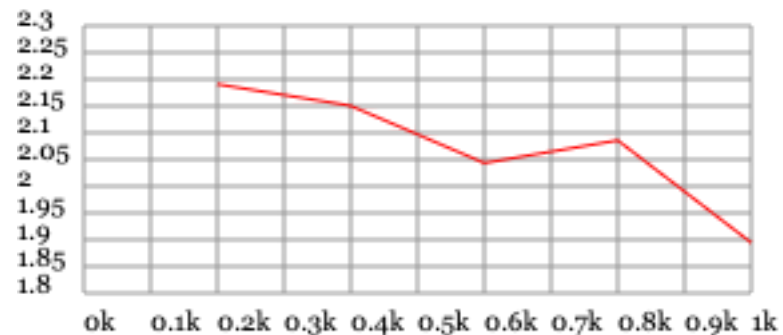
Learning rate:

Momentum:

Batch size:

Weight decay:

Loss:



Instantiate a Network and Trainer

```

layer_defs = [];
layer_defs.push({type:'input', out_sx:32, out_sy:32, out_depth:3});
layer_defs.push({type:'conv', sx:5, filters:16, stride:1, pad:2, activation:'relu'});
layer_defs.push({type:'pool', sx:2, stride:2});
layer_defs.push({type:'conv', sx:5, filters:20, stride:1, pad:2, activation:'relu'});
layer_defs.push({type:'pool', sx:2, stride:2});
layer_defs.push({type:'conv', sx:5, filters:20, stride:1, pad:2, activation:'relu'});
layer_defs.push({type:'pool', sx:2, stride:2});
layer_defs.push({type:'softmax', num_classes:10});

net = new convnetjs.Net();
net.makeLayers(layer_defs);

trainer = new convnetjs.SGDTrainer(net, {method:'adadelata', batch_size:4, l2_decay:0.0001});


```

Network Visualization

input (32x32x3)

max activation: 0.48823, min: -0.47648
max gradient: 0.00996, min: -0.00899

Activations:




conv (32x32x16)

filter size 5x5x3, stride 1


max activation: 1.30636, min: -1.0832
max gradient: 0.00579, min: -0.00704
parameters: $16 \times 5 \times 5 \times 3 + 16 = 1216$

Activations:



Activation Gradients:





Weights:



Weight Gradients:




relu (32x32x16)


max activation: 1.30636, min: 0

max gradient: 0.00777, min: -0.00704

Activations:



Activation Gradients:




pool (16x16x16)
 pooling size 2x2, stride 2
 max activation: 1.30636, min: 0
 max gradient: 0.00777, min: -0.00704

Activations:




Activation Gradients:



conv (16x16x20)
 filter size 5x5x16, stride 1
 max activation: 3.23825, min: -4.47611
 max gradient: 0.01436, min: -0.01682
 parameters: $20 \times 5 \times 5 \times 16 + 20 = 8020$

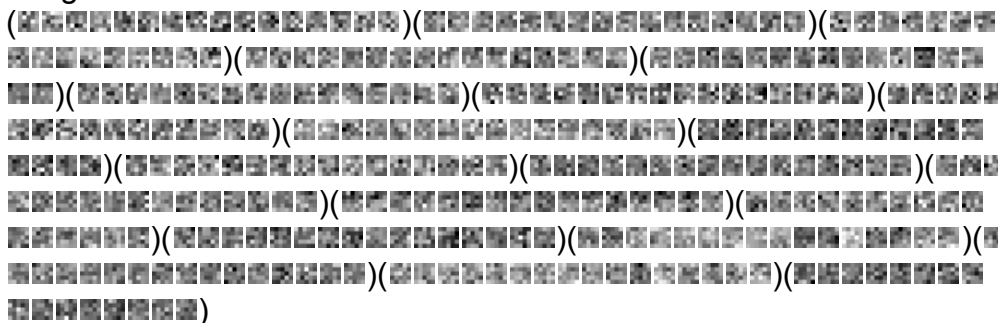
Activations:



Activation Gradients:



Weights:

()

Weight Gradients:

()



relu (16x16x20)


max activation: 3.23825, min: 0

max gradient: 0.01436, min: -0.01682

Activations:



Activation Gradients:



pool (8x8x20)

pooling size 2x2, stride 2

max activation: 3.23825, min: 0

max gradient: 0.01436, min: -0.01682

Activations:



Activation Gradients:



conv (8x8x20)


filter size 5x5x20, stride 1

max activation: 3.34058, min: -4.87992

max gradient: 0.0347, min: -0.03621

parameters: $20 \times 5 \times 5 \times 20 + 20 = 10020$


Activations:




Activation Gradients:




Weights:



Weight Gradients:





relu (8x8x20)

max activation: 3.34058, min: 0

max gradient: 0.05703, min: -0.03851

Activations:



Activation Gradients:



pool (4x4x20)

pooling size 2x2, stride 2

max activation: 3.34058, min: 0

max gradient: 0.05703, min: -0.03851

Activations:



Activation Gradients:



fc (1x1x10)

max activation: 1.69487, min: -8.53016

max gradient: 0.10433, min: -0.19265

parameters: $10 \times 320 + 10 = 3210$

Activations:



Activation Gradients:



softmax (1x1x10)

max activation: 0.80735, min: 0.00002

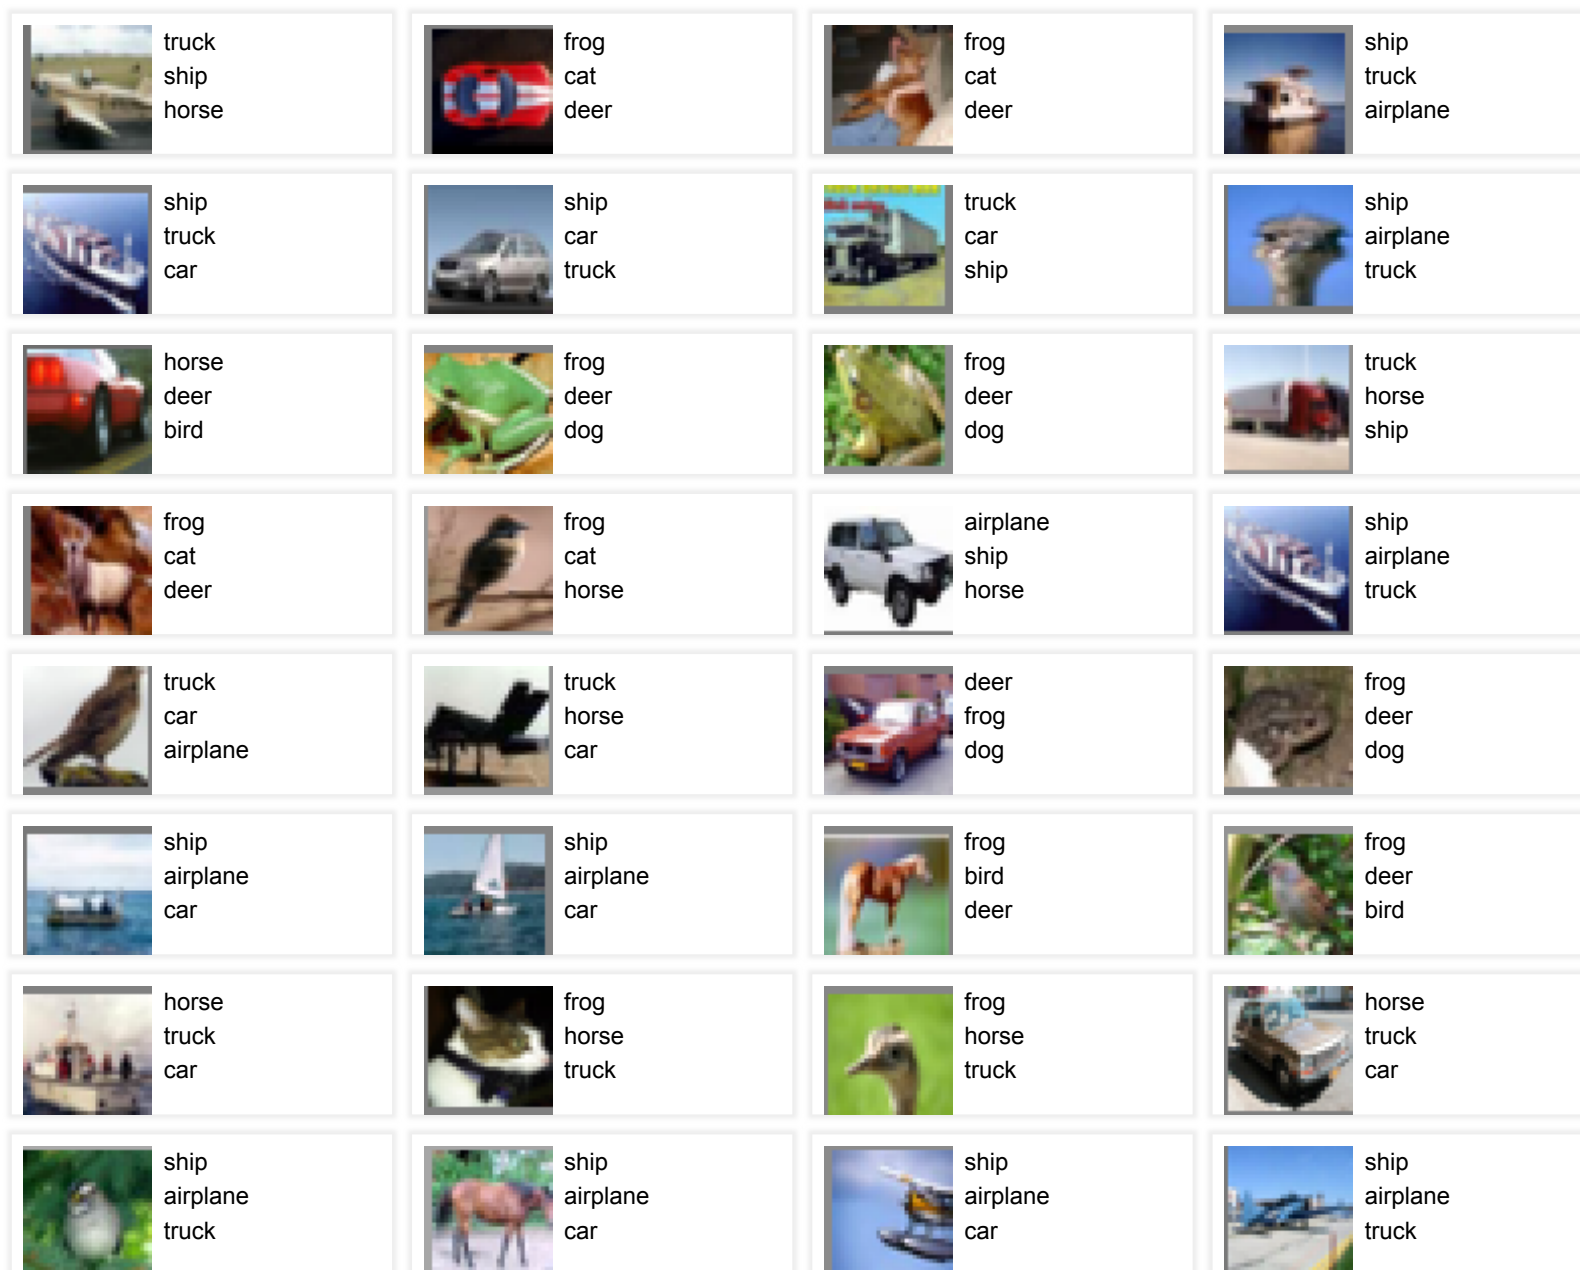
max gradient: 0, min: 0


Activations:



Example predictions on Test set

test accuracy based on last 200 test images: 0.275





deer
frog
cat




airplane
car
ship



deer
frog
cat



deer
truck
frog




cat
deer
dog



airplane
car
ship



ship
airplane
car



car
airplane
cat