

ConvNetJS MNIST demo

Description

This demo trains a Convolutional Neural Network on the [MNIST digits dataset](#) in your browser, with nothing but Javascript. The dataset is fairly easy and one should expect to get somewhere around 99% accuracy within few minutes. I used [this python script](#) to parse the [original files](#) into batches of images that can be easily loaded into page DOM with img tags.

This network takes a 28x28 MNIST image and crops a random 24x24 window before training on it (this technique is called data augmentation and improves generalization). Similarly to do prediction, 4 random crops are sampled and the probabilities across all crops are averaged to produce final predictions. The network runs at about 5ms for both forward and backward pass on my reasonably decent Ubuntu+Chrome machine.

By default, in this demo we're using Adadelta which is one of per-parameter adaptive step size methods, so we don't have to worry about changing learning rates or momentum over time. However, I still included the text fields for changing these if you'd like to play around with SGD+Momentum trainer.


Report questions/bugs/suggestions to [@karpathy](#).

Training Stats

pause

Forward time per example: 2ms
Backprop time per example: 2ms
Classification loss: 0.50576
L2 Weight decay loss: 0.00135
Training accuracy: 0.87
Validation accuracy: 0.78
Examples seen: 2000
Learning rate:
Momentum:
Batch size:
Weight decay:

Loss:



Instantiate a Network and Trainer

```
layer_defs = [];
layer_defs.push({type:'input', out_sx:24, out_sy:24, out_depth:1});
layer_defs.push({type:'conv', sx:5, filters:8, stride:1, pad:2, activation:'relu'});
layer_defs.push({type:'pool', sx:2, stride:2});
layer_defs.push({type:'conv', sx:5, filters:16, stride:1, pad:2, activation:'relu'});
layer_defs.push({type:'pool', sx:3, stride:3});
layer_defs.push({type:'softmax', num_classes:10});

net = new convnetjs.Net();
net.makeLayers(layer_defs);

trainer = new convnetjs.SGDTrainer(net, {method:'adadelata', batch_size:20, l2_decay:0.001});
```

change network

Network Visualization

input (24x24x1)
max activation: 1, min: 0
max gradient: 0.46304, min: -0.3949

Activations:



Activation Gradients:



```
conv (24x24x8)
filter size 5x5x1, stride 1
max activation: 2.47644, min: -1.43224
max gradient: 0.23071, min: -0.24962
parameters: 8x5x5x1+8 = 208
```


Activations:



Activation Gradients:



Weights:



Weight Gradients:




relu (24x24x8)
max activation: 2.47644, min: 0
max gradient: 0.23071, min: -0.24962

Activations:



Activation Gradients:




pool (12x12x8)
pooling size 2x2, stride 2
max activation: 2.47644, min: 0
max gradient: 0.23071, min: -0.24962

Activations:



Activation Gradients:



```
conv (12x12x16)
filter size 5x5x8, stride 1
max activation: 4.07014, min: -8.45024
max gradient: 0.28477, min: -0.39567
parameters: 16x5x5x8+16 = 3216
```

Activations:



Activation Gradients:




Weights:

[illegible]


Weight Gradients:

relu (12x12x16)
max activation: 4.07014, min: 0
max gradient: 0.28477, min: -0.39567

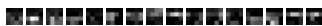

Activations:





Activation Gradients:



pool (4x4x16)
 pooling size 3x3, stride 3
 max activation: 4.07014, min: 0
 max gradient: 0.28477, min: -0.39567

Activations:

 Activation Gradients:


fc (1x1x10)
 max activation: 2.31402, min: -5.56164
 max gradient: 0.75314, min: -0.89211
 parameters: 10x256+10 = 2570

Activations:

 Activation Gradients:


softmax (1x1x10)
 max activation: 0.75314, min: 0.00028
 max gradient: 0, min: 0

Activations:


Example predictions on Test set

